

Class06: R Functions

Julia Di Silvestri (PID: A16950824)

2024-01-25

R Functions

Functions are how we get stuff done. We call functions to do everything useful in R.

One cool thing about R is that it makes writing your own functions comparatively easy.

All functions in R have at least three things:

- A **name** (we get to pick this)
- One or more **input arguments** (the input to our function)
- The **body** (lines of code that do the work)

#| eval: false will allow invalid code to render by echoing it

```
funname <- function(input1, input2) {  
  #The body with R code  
}
```

Let's write a silly first function to add two numbers:

```
x <- 5  
y <- 1  
x + y
```

```
[1] 6
```

```
addme <- function(x, y) { x + y }
```

```
addme(279, 5678)
```

```
[1] 5957
```

To assign one element with a default:

```
addme <- function(x, y=1) {x + y}
```

```
addme(10)
```

```
[1] 11
```

Lab for Today

Question 1

Writing a “grade” function

First, we assign vectors to each student with each of their grades:

```
student1 <- c(100, 100, 100, 100, 100, 100, 100, 90)
student2 <- c(100, NA, 90, 90, 90, 90, 97, 80)
student3 <- c(90, NA, NA, NA, NA, NA, NA, NA)
```

Next, we will start to form the function. Step one is to get the function to identify a student's lowest score:

```
which.min(student1)
```

```
[1] 8
```

```
which.min(student2)
```

```
[1] 8
```

```
which.min(student3)
```

```
[1] 1
```

Next, we will find the average of each student's scores:

```
mean(student1)
```

```
[1] 98.75
```

```
mean(student2, na.rm = T)
```

```
[1] 91
```

```
mean(student3, na.rm = T)
```

```
[1] 90
```

This is currently not fair – student 3 should not have a mean of 90.

We will move on for now. Things worked for student 1. Now we want to drop the lowest score before getting the `mean()`. Using `vector[-x]` will spit out every value in that vector except for the xth value.

```
# Find lowest score  
which.min(student1)
```

```
[1] 8
```

```
# Remove lowest score  
student1[-8]
```

```
[1] 100 100 100 100 100 100 100
```

Now to put it together:

```
# Find mean with lowest score removed  
mean(student1[-which.min(student1)])
```

```
[1] 100
```

Nice it worked!

A common shortcut and use `x` as my input so that I don't have to keep typing `student1`

```
x <- student1
mean( x[ -which.min(x)])
```

```
[1] 100
```

We still have the problem of missing values.

One idea is to replace NA values with zero.

```
y <- 1:5
y[y == 3] <- 10000
y
```

```
[1]      1      2 10000      4      5
```

^ This method will not work for NAs because there is not data to change in the case of an NA

To find location if NA in a vector:

```
is.na(student2)
```

```
[1] FALSE  TRUE FALSE FALSE FALSE FALSE FALSE
```

How to set NA to 0:

```
b <- student2
b[is.na(b)] <- 0
mean( b[ -which.min(b)])
```

```
[1] 91
```

Now writing a function with this:

```
grade <- function(x) {
  #change NA to 0
  x[is.na(x)] <- 0
  #drop lowest score and average
  mean( x[ -which.min(x)])}
```

```
grade(student1)
```

```
[1] 100
```

```
grade(student2)
```

```
[1] 91
```

```
grade(student3)
```

```
[1] 12.85714
```

To import CSV file:

```
url <- "https://tinyurl.com/gradeinput"  
gradebook <- read.csv(url, row.names = 1)  
gradebook
```

	hw1	hw2	hw3	hw4	hw5
student-1	100	73	100	88	79
student-2	85	64	78	89	78
student-3	83	69	77	100	77
student-4	88	NA	73	100	76
student-5	88	100	75	86	79
student-6	89	78	100	89	77
student-7	89	100	74	87	100
student-8	89	100	76	86	100
student-9	86	100	77	88	77
student-10	89	72	79	NA	76
student-11	82	66	78	84	100
student-12	100	70	75	92	100
student-13	89	100	76	100	80
student-14	85	100	77	89	76
student-15	85	65	76	89	NA
student-16	92	100	74	89	77
student-17	88	63	100	86	78
student-18	91	NA	100	87	100
student-19	91	68	75	86	79
student-20	91	68	76	88	76

To apply our function to “gradebook”, we use `apply()`:

```
# apply function works like: apply(dataframe/matrix, 1 (row) or 2 (col) or c(1,2) (both),
results <- apply(gradebook, 1, grade)
results
```

student-1	student-2	student-3	student-4	student-5	student-6	student-7
91.75	82.50	84.25	84.25	88.25	89.00	94.00
student-8	student-9	student-10	student-11	student-12	student-13	student-14
93.75	87.75	79.00	86.00	91.75	92.25	87.75
student-15	student-16	student-17	student-18	student-19	student-20	
78.75	89.50	88.00	94.50	82.75	82.75	

Question 2

To pull out the top scoring student, use `which.max()`

```
#who scored the highest
which.max(results)
```

```
student-18
18
```

```
#what did they score
max(results)
```

```
[1] 94.5
```

Question 3

To determine the homework averages across the whole class, we will apply a function over the columns instead of the rows

```
apply(gradebook, 2, mean)
```

hw1	hw2	hw3	hw4	hw5
89.0	NA	80.8	NA	NA

This did not work because we need to write a new function that drops the NAs again

```
hwavgs <- apply(gradebook, 2, mean, na.rm = T)
hwavgs
```

```
      hw1      hw2      hw3      hw4      hw5
89.00000 80.88889 80.80000 89.63158 83.42105
```

Now we can pull out the minimum from these answers:

```
# which homework has the lowest average
which.min(hwavgs)
```

```
hw3
3
```

```
# what was the average
min(hwavgs)
```

```
[1] 80.8
```

Now we will use another method to try to weed out biases in the average

```
hwsums <- apply(gradebook, 2, sum, na.rm = T)
which.min(hwsums)
```

```
hw2
2
```

This answer indicates that there was an outlier in HW3 that was skewing the average. Homework two was more consistently the low-scoring.

Question 4

```
# make all (or mask) NAs to zero
mask <- gradebook
mask[is.na(mask)] <- 0
```

We can use the `cor()` function for correlation analysis

```
cor(mask$hw5, results)
```

```
[1] 0.6325982
```

```
cor(mask$hw3, results)
```

```
[1] 0.3042561
```

Homework 5 is much more correlated. How do we apply this across the whole gradebook?

```
cors <- apply(mask, 2, cor, results)
cors
```

	hw1	hw2	hw3	hw4	hw5
	0.4250204	0.1767780	0.3042561	0.3810884	0.6325982

We can see that hw5 is the most correlated, and that maybe hw2 should be reconsidered.